

Credit Card Fraud

May 3, 2018

1 Credit Card Fraud

1.0.1 About Dataset

The datasets contains transactions made by credit cards in September 2013 by european cardholders. This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.

1.0.2 Dataset at a Glance

```
In [1]: %matplotlib inline
```

```
import itertools
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.gridspec as gridspec
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_validate
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import average_precision_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from imblearn.over_sampling import SMOTE
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Dropout

pd.set_option("display.max_columns", 31)

transactions = pd.read_csv("creditcard.csv")
transactions.describe()
```

```

/home/firabby/pyML/lib/python3.6/site-packages/h5py/__init__.py:36: FutureWarning: Conversion
  from ._conv import register_converters as _register_converters
Using TensorFlow backend.

```

```

Out[1]:

```

	Time	V1	V2	V3	V4 \
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.165980e-15	3.416908e-16	-1.373150e-15	2.086869e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01

	V5	V6	V7	V8	V9 \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	9.604066e-16	1.490107e-15	-5.556467e-16	1.177556e-16	-2.406455e-15
std	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00
min	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+01
25%	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01
50%	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02
75%	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01
max	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+01

	V10	V11	V12	V13	V14 \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	2.239751e-15	1.673327e-15	-1.254995e-15	8.176030e-16	1.206296e-15
std	1.088850e+00	1.020713e+00	9.992014e-01	9.952742e-01	9.585956e-01
min	-2.458826e+01	-4.797473e+00	-1.868371e+01	-5.791881e+00	-1.921433e+01
25%	-5.354257e-01	-7.624942e-01	-4.055715e-01	-6.485393e-01	-4.255740e-01
50%	-9.291738e-02	-3.275735e-02	1.400326e-01	-1.356806e-02	5.060132e-02
75%	4.539234e-01	7.395934e-01	6.182380e-01	6.625050e-01	4.931498e-01
max	2.374514e+01	1.201891e+01	7.848392e+00	7.126883e+00	1.052677e+01

	V15	V16	V17	V18	V19 \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	4.913003e-15	1.437666e-15	-3.800113e-16	9.572133e-16	1.039817e-15
std	9.153160e-01	8.762529e-01	8.493371e-01	8.381762e-01	8.140405e-01
min	-4.498945e+00	-1.412985e+01	-2.516280e+01	-9.498746e+00	-7.213527e+00
25%	-5.828843e-01	-4.680368e-01	-4.837483e-01	-4.988498e-01	-4.562989e-01
50%	4.807155e-02	6.641332e-02	-6.567575e-02	-3.636312e-03	3.734823e-03
75%	6.488208e-01	5.232963e-01	3.996750e-01	5.008067e-01	4.589494e-01
max	8.877742e+00	1.731511e+01	9.253526e+00	5.041069e+00	5.591971e+00

	V20	V21	V22	V23	V24 \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	6.406703e-16	1.656562e-16	-3.444850e-16	2.578648e-16	4.471968e-15

std	7.709250e-01	7.345240e-01	7.257016e-01	6.244603e-01	6.056471e-01
min	-5.449772e+01	-3.483038e+01	-1.093314e+01	-4.480774e+01	-2.836627e+00
25%	-2.117214e-01	-2.283949e-01	-5.423504e-01	-1.618463e-01	-3.545861e-01
50%	-6.248109e-02	-2.945017e-02	6.781943e-03	-1.119293e-02	4.097606e-02
75%	1.330408e-01	1.863772e-01	5.285536e-01	1.476421e-01	4.395266e-01
max	3.942090e+01	2.720284e+01	1.050309e+01	2.252841e+01	4.584549e+00

	V25	V26	V27	V28	Amount \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	284807.000000
mean	5.340915e-16	1.687098e-15	-3.666453e-16	-1.220404e-16	88.349619
std	5.212781e-01	4.822270e-01	4.036325e-01	3.300833e-01	250.120109
min	-1.029540e+01	-2.604551e+00	-2.256568e+01	-1.543008e+01	0.000000
25%	-3.171451e-01	-3.269839e-01	-7.083953e-02	-5.295979e-02	5.600000
50%	1.659350e-02	-5.213911e-02	1.342146e-03	1.124383e-02	22.000000
75%	3.507156e-01	2.409522e-01	9.104512e-02	7.827995e-02	77.165000
max	7.519589e+00	3.517346e+00	3.161220e+01	3.384781e+01	25691.160000

	Class
count	284807.000000
mean	0.001727
std	0.041527
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

In [2]: transactions.head()

Out [2]:	Time	V1	V2	V3	V4	V5	V6	V7 \
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941

	V8	V9	V10	V11	V12	V13	V14 \
0	0.098698	0.363787	0.090794	-0.551600	-0.617801	-0.991390	-0.311169
1	0.085102	-0.255425	-0.166974	1.612727	1.065235	0.489095	-0.143772
2	0.247676	-1.514654	0.207643	0.624501	0.066084	0.717293	-0.165946
3	0.377436	-1.387024	-0.054952	-0.226487	0.178228	0.507757	-0.287924
4	-0.270533	0.817739	0.753074	-0.822843	0.538196	1.345852	-1.119670

	V15	V16	V17	V18	V19	V20	V21 \
0	1.468177	-0.470401	0.207971	0.025791	0.403993	0.251412	-0.018307
1	0.635558	0.463917	-0.114805	-0.183361	-0.145783	-0.069083	-0.225775
2	2.345865	-2.890083	1.109969	-0.121359	-2.261857	0.524980	0.247998
3	-0.631418	-1.059647	-0.684093	1.965775	-1.232622	-0.208038	-0.108300

```
4  0.175121 -0.451449 -0.237033 -0.038195  0.803487  0.408542 -0.009431
```

```

      V22      V23      V24      V25      V26      V27      V28  \
0  0.277838 -0.110474  0.066928  0.128539 -0.189115  0.133558 -0.021053
1 -0.638672  0.101288 -0.339846  0.167170  0.125895 -0.008983  0.014724
2  0.771679  0.909412 -0.689281 -0.327642 -0.139097 -0.055353 -0.059752
3  0.005274 -0.190321 -1.175575  0.647376 -0.221929  0.062723  0.061458
4  0.798278 -0.137458  0.141267 -0.206010  0.502292  0.219422  0.215153

```

```

      Amount  Class
0  149.62      0
1   2.69      0
2  378.66      0
3  123.50      0
4   69.99      0

```

```
In [3]: # check if there is any missing value
transactions.isnull().sum()
```


```
Out[3]: Time      0
V1          0
V2          0
V3          0
V4          0
V5          0
V6          0
V7          0
V8          0
V9          0
V10         0
V11         0
V12         0
V13         0
V14         0
V15         0
V16         0
V17         0
V18         0
V19         0
V20         0
V21         0
V22         0
V23         0
V24         0
V25         0
V26         0
V27         0
V28         0
```

```
Amount    0
Class     0
dtype: int64
```

```
In [4]: transactions_count = pd.value_counts(transactions.Class)
figure, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 7))
ax1.bar(transactions_count.index, transactions_count.values)
plt.title("Bar Chart")
plt.xlabel("Class")
plt.ylabel("Frequency")

ax2.pie(transactions_count)
plt.title("Pie Chart")
```

```
Out[4]: Text(0.5,1,'Pie Chart')
```




From these charts it is clear that the data is very imbalanced.
Let's see how time compares to both fraud and normal transactions.

```
In [5]: figure, (ax1, ax2) = plt.subplots(2, 1, sharex=True, figsize=(15, 15))

ax1.hist(transactions.Time[transactions.Class == 1])
ax1.set_title('Fraud')

ax2.hist(transactions.Time[transactions.Class == 0])
ax2.set_title('Normal')

plt.xlabel('Time (in Seconds)')
plt.ylabel('Number of Transactions')
plt.show()
```



We can see that fraud transaction were still being made a little after the peak hour. This information can help us detect fraud transaction during off-peak hour.

Let us see if we can find any relation in amount for fraud and normal transaction.


In [6]: `figure, (ax1, ax2) = plt.subplots(2, 1, sharex=True, figsize=(15, 15))`

```
ax1.hist(transactions.Amount[transactions.Class == 1])
ax1.set_title('Fraud')
```

```
ax2.hist(transactions.Amount[transactions.Class == 0])
ax2.set_title('Normal')
```

```
plt.xlabel('Amount')
plt.ylabel('Number of Transactions')
plt.yscale('log')
```

```
plt.show()
```



We can see that fraud transactions were made in smaller amount.

1.0.3 Normalize the Data

Only the amounts are not normalized. So, we need to normalize the amounts.

```
In [7]: transactions["normalizedAmount"] = StandardScaler().fit_transform(transactions.Amount.values)
```

Drop time and amount column, as we no longer need them.

```
In [8]: transactions = transactions.drop(["Amount", "Time"], axis=1)
```

Let's look at our new dataset.

In [9]: transactions.describe()

```
Out [9]:
```

	V1	V2	V3	V4	V5	\
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	
mean	1.165980e-15	3.416908e-16	-1.373150e-15	2.086869e-15	9.604066e-16	
std	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	
min	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	
25%	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	
50%	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	
75%	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	
max	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	

	V6	V7	V8	V9	V10	\
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	
mean	1.490107e-15	-5.556467e-16	1.177556e-16	-2.406455e-15	2.239751e-15	
std	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00	1.088850e+00	
min	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+01	-2.458826e+01	
25%	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01	-5.354257e-01	
50%	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02	-9.291738e-02	
75%	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01	4.539234e-01	
max	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+01	2.374514e+01	

	V11	V12	V13	V14	V15	\
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	
mean	1.673327e-15	-1.254995e-15	8.176030e-16	1.206296e-15	4.913003e-15	
std	1.020713e+00	9.992014e-01	9.952742e-01	9.585956e-01	9.153160e-01	
min	-4.797473e+00	-1.868371e+01	-5.791881e+00	-1.921433e+01	-4.498945e+00	
25%	-7.624942e-01	-4.055715e-01	-6.485393e-01	-4.255740e-01	-5.828843e-01	
50%	-3.275735e-02	1.400326e-01	-1.356806e-02	5.060132e-02	4.807155e-02	
75%	7.395934e-01	6.182380e-01	6.625050e-01	4.931498e-01	6.488208e-01	
max	1.201891e+01	7.848392e+00	7.126883e+00	1.052677e+01	8.877742e+00	

	V16	V17	V18	V19	V20	\
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	
mean	1.437666e-15	-3.800113e-16	9.572133e-16	1.039817e-15	6.406703e-16	
std	8.762529e-01	8.493371e-01	8.381762e-01	8.140405e-01	7.709250e-01	
min	-1.412985e+01	-2.516280e+01	-9.498746e+00	-7.213527e+00	-5.449772e+01	
25%	-4.680368e-01	-4.837483e-01	-4.988498e-01	-4.562989e-01	-2.117214e-01	
50%	6.641332e-02	-6.567575e-02	-3.636312e-03	3.734823e-03	-6.248109e-02	
75%	5.232963e-01	3.996750e-01	5.008067e-01	4.589494e-01	1.330408e-01	
max	1.731511e+01	9.253526e+00	5.041069e+00	5.591971e+00	3.942090e+01	

	V21	V22	V23	V24	V25	\
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	
mean	1.656562e-16	-3.444850e-16	2.578648e-16	4.471968e-15	5.340915e-16	
std	7.345240e-01	7.257016e-01	6.244603e-01	6.056471e-01	5.212781e-01	
min	-3.483038e+01	-1.093314e+01	-4.480774e+01	-2.836627e+00	-1.029540e+01	
25%	-2.283949e-01	-5.423504e-01	-1.618463e-01	-3.545861e-01	-3.171451e-01	

50%	-2.945017e-02	6.781943e-03	-1.119293e-02	4.097606e-02	1.659350e-02
75%	1.863772e-01	5.285536e-01	1.476421e-01	4.395266e-01	3.507156e-01
max	2.720284e+01	1.050309e+01	2.252841e+01	4.584549e+00	7.519589e+00

	V26	V27	V28	Class \
count	2.848070e+05	2.848070e+05	2.848070e+05	284807.000000
mean	1.687098e-15	-3.666453e-16	-1.220404e-16	0.001727
std	4.822270e-01	4.036325e-01	3.300833e-01	0.041527
min	-2.604551e+00	-2.256568e+01	-1.543008e+01	0.000000
25%	-3.269839e-01	-7.083953e-02	-5.295979e-02	0.000000
50%	-5.213911e-02	1.342146e-03	1.124383e-02	0.000000
75%	2.409522e-01	9.104512e-02	7.827995e-02	0.000000
max	3.517346e+00	3.161220e+01	3.384781e+01	1.000000

	normalizedAmount
count	2.848070e+05
mean	2.913952e-17
std	1.000002e+00
min	-3.532294e-01
25%	-3.308401e-01
50%	-2.652715e-01
75%	-4.471707e-02
max	1.023622e+02

In [10]: transactions.head()

Out [10]:

	V1	V2	V3	V4	V5	V6	V7 \
0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599
1	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803
2	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461
3	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609
4	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941

	V8	V9	V10	V11	V12	V13	V14 \
0	0.098698	0.363787	0.090794	-0.551600	-0.617801	-0.991390	-0.311169
1	0.085102	-0.255425	-0.166974	1.612727	1.065235	0.489095	-0.143772
2	0.247676	-1.514654	0.207643	0.624501	0.066084	0.717293	-0.165946
3	0.377436	-1.387024	-0.054952	-0.226487	0.178228	0.507757	-0.287924
4	-0.270533	0.817739	0.753074	-0.822843	0.538196	1.345852	-1.119670

	V15	V16	V17	V18	V19	V20	V21 \
0	1.468177	-0.470401	0.207971	0.025791	0.403993	0.251412	-0.018307
1	0.635558	0.463917	-0.114805	-0.183361	-0.145783	-0.069083	-0.225775
2	2.345865	-2.890083	1.109969	-0.121359	-2.261857	0.524980	0.247998
3	-0.631418	-1.059647	-0.684093	1.965775	-1.232622	-0.208038	-0.108300
4	0.175121	-0.451449	-0.237033	-0.038195	0.803487	0.408542	-0.009431

	V22	V23	V24	V25	V26	V27	V28 \
--	-----	-----	-----	-----	-----	-----	-------

```

0  0.277838 -0.110474  0.066928  0.128539 -0.189115  0.133558 -0.021053
1 -0.638672  0.101288 -0.339846  0.167170  0.125895 -0.008983  0.014724
2  0.771679  0.909412 -0.689281 -0.327642 -0.139097 -0.055353 -0.059752
3  0.005274 -0.190321 -1.175575  0.647376 -0.221929  0.062723  0.061458
4  0.798278 -0.137458  0.141267 -0.206010  0.502292  0.219422  0.215153

```

```

      Class  normalizedAmount
0         0          0.244964
1         0         -0.342475
2         0          1.160686
3         0          0.140534
4         0         -0.073403

```

1.0.4 Splitting Data for Train and Test Set with and without Resampling

Data without resampling

```

In [11]: X = transactions.iloc[:, transactions.columns != 'Class'].values
        y = transactions.iloc[:, transactions.columns == 'Class'].values

```

Undersampling Normal Transactions

```

In [12]: fraud_transactions_count = len(transactions[transactions.Class == 1])

        fraud_indices = np.array(transactions[transactions.Class == 1].index)
        normal_indices = transactions[transactions.Class == 0].index

        random_normal_indices = np.random.choice(normal_indices, fraud_transactions_count, replace=True)

        undersampled_indices = np.concatenate([fraud_indices, random_normal_indices])

        undersampled_transactions = transactions.iloc[undersampled_indices, :]

        X_undersampled = undersampled_transactions.iloc[:, undersampled_transactions.columns != 'Class']
        y_undersampled = undersampled_transactions.iloc[:, undersampled_transactions.columns == 'Class']

        print("Normal transactions: ", len(undersampled_transactions[undersampled_transactions.Class == 0]))
        print("Fraud transactions: ", len(undersampled_transactions[undersampled_transactions.Class == 1]))
        print("Total transactions in resampled data: ", len(undersampled_transactions))

```

```

Normal transactions:  0.5
Fraud transactions:  0.5
Total transactions in resampled data:  984

```

Oversampling using SMOTE

```

In [13]: sm = SMOTE(kind="regular")

In [14]: X_oversampled, y_oversampled = sm.fit_sample(X, y.ravel())

In [15]: oversampled_transactions = np.hstack((X_oversampled, y_oversampled.reshape(-1, 1)))
         print("Normal transactions: ", len(oversampled_transactions[oversampled_transactions[:, 1] == 0]))
         print("Fraud transactions: ", len(oversampled_transactions[oversampled_transactions[:, 1] == 1]))
         print("Total transactions in resampled data: ", len(oversampled_transactions))

Normal transactions:  0.5
Fraud transactions:  0.5
Total transactions in resampled data:  568630

```

1.0.5 Splitting into Train and Test Data

```

In [16]: X_train, X_test, y_train, y_test = train_test_split(X, y.ravel(), test_size=0.2, random_state=42)
         X_undersampled_train, X_undersampled_test, y_undersampled_train, y_undersampled_test = train_test_split(X_oversampled_train, X_oversampled_test, y_oversampled_train, y_oversampled_test, test_size=0.2, random_state=42)

In [17]: def get_cv_scores(clf, x, y):
         scoring = ['precision_micro', 'recall_micro', 'precision_macro', 'recall_macro']
         scores = cross_validate(clf, x, y, scoring=scoring, cv=5, return_train_score=False)
         scores_mean = {}
         print("Average Training Time: ", scores["fit_time"].mean(), "(+/- %0.2f)" % (scores["fit_time"].std()))
         print("Average Training Time: ", scores["score_time"].mean(), "(+/- %0.2f)" % (scores["score_time"].std()))
         print("Average Test Precision Macro: ", scores["test_precision_macro"].mean(), "(+/- %0.2f)" % (scores["test_precision_macro"].std()))
         print("Average Test Precision Micro: ", scores["test_precision_micro"].mean(), "(+/- %0.2f)" % (scores["test_precision_micro"].std()))
         print("Average Test Recall Macro: ", scores["test_recall_macro"].mean(), "(+/- %0.2f)" % (scores["test_recall_macro"].std()))
         print("Average Test Recall Micro: ", scores["test_recall_micro"].mean(), "(+/- %0.2f)" % (scores["test_recall_micro"].std()))

         scores_mean["test_precision_macro"] = scores["test_precision_macro"].mean()
         scores_mean["test_precision_micro"] = scores["test_precision_micro"].mean()
         scores_mean["test_recall_macro"] = scores["test_recall_macro"].mean()
         scores_mean["test_recall_micro"] = scores["test_recall_micro"].mean()

         return scores_mean

clf = LogisticRegression(C=0.01, random_state=42)

print("Original Data")
print("=====\\n")
org_lr_score = get_cv_scores(clf, X_train, y_train)
print()

print("Undersampled Data")
print("=====\\n")
und_lr_score = get_cv_scores(clf, X_undersampled_train, y_undersampled_train)

```

```

print()

print("Oversampled Data")
print("=====\\n")
ovr_lr_score = get_cv_scores(clf, X_oversampled_train, y_oversampled_train)

```

Original Data

=====

```

Average Training Time:  1.173475170135498 (+/- 0.10)
Average Training Time:  0.04214701652526855 (+/- 0.01)
Average Test Precision Macro:  0.9458818378730246 (+/- 0.01)
Average Test Precision Micro:  0.9991573221523449 (+/- 0.00)
Average Test Recall Macro:  0.7919377991122797 (+/- 0.03)
Average Test Recall Micro:  0.9991573221523449 (+/- 0.00)

```

Undersampled Data

=====

```

Average Training Time:  0.0049957275390625 (+/- 0.00)
Average Training Time:  0.002244901657104492 (+/- 0.00)
Average Test Precision Macro:  0.9192994880497036 (+/- 0.01)
Average Test Precision Micro:  0.9173814441088488 (+/- 0.01)
Average Test Recall Macro:  0.9173482635507952 (+/- 0.01)
Average Test Recall Micro:  0.9173814441088488 (+/- 0.01)

```

Oversampled Data

=====

```

Average Training Time:  4.735983037948609 (+/- 0.22)
Average Training Time:  0.09955897331237792 (+/- 0.00)
Average Test Precision Macro:  0.9467523161299045 (+/- 0.00)
Average Test Precision Micro:  0.9450763222483586 (+/- 0.00)
Average Test Recall Macro:  0.9450611497209722 (+/- 0.00)
Average Test Recall Micro:  0.9450763222483586 (+/- 0.00)

```

In [18]: clf = RandomForestClassifier(random_state=42)

```

print("Original Data")
print("=====\\n")
org_rf_score = get_cv_scores(clf, X_train, y_train)
print()

print("Undersampled Data")
print("=====\\n")
und_rf_score = get_cv_scores(clf, X_undersampled_train, y_undersampled_train)
print()

```

```

print("Oversampled Data")
print("=====\n")
ovr_rf_score = get_cv_scores(clf, X_oversampled_train, y_oversampled_train)

```

Original Data

=====

```

Average Training Time:  14.825970077514649 (+/- 1.42)
Average Training Time:  0.17570128440856933 (+/- 0.01)
Average Test Precision Macro:  0.9721289734137869 (+/- 0.02)
Average Test Precision Micro:  0.9994952711967147 (+/- 0.00)
Average Test Recall Macro:  0.8769646505148264 (+/- 0.03)
Average Test Recall Micro:  0.9994952711967147 (+/- 0.00)

```

Undersampled Data

=====

```

Average Training Time:  0.01947441101074219 (+/- 0.00)
Average Training Time:  0.004582881927490234 (+/- 0.00)
Average Test Precision Macro:  0.927227247385666 (+/- 0.01)
Average Test Precision Micro:  0.9237429116611573 (+/- 0.01)
Average Test Recall Macro:  0.9236773774748459 (+/- 0.01)
Average Test Recall Micro:  0.9237429116611573 (+/- 0.01)

```

Oversampled Data

=====

```

Average Training Time:  27.69242305755615 (+/- 0.86)
Average Training Time:  0.47842750549316404 (+/- 0.01)
Average Test Precision Macro:  0.9998175263441713 (+/- 0.00)
Average Test Precision Micro:  0.9998175438875245 (+/- 0.00)
Average Test Recall Macro:  0.9998175757121663 (+/- 0.00)
Average Test Recall Micro:  0.9998175438875245 (+/- 0.00)

```


Random forest takes longer time to train as data increases.

```

In [22]: plt.figure(figsize=(15, 10))
         plt.bar(range(len(ovr_rf_score)), list(ovr_rf_score.values()), align='center', width=0.8)
         plt.xticks(range(len(ovr_rf_score)), list(ovr_rf_score.keys()))

         plt.bar(np.arange(len(ovr_lr_score))+0.2, list(ovr_lr_score.values()), align='center', width=0.8)
         plt.xticks(range(len(ovr_lr_score)), list(ovr_lr_score.keys()))
         plt.legend(["Random Forest", "Linear Regression"])
         plt.title("Precision, Recall Comparison")
         plt.show()

```



From the above experiments, we can see that oversampled data with Synthetic Minority Over-sampling Technique, we get the best precision and recall using Random Forest Classifier. Let's train the classifier and evaluate the model.

```
In [20]: clf.fit(X_oversampled_train, y_oversampled_train)
pred = clf.predict(X_oversampled_test)
class_names = ["Normal", "Fraud"]
report = classification_report(y_oversampled_test, pred, target_names = class_names)
print(report)
print("Accuracy Score: ", accuracy_score(y_oversampled_test, pred))

average_precision = average_precision_score(y_oversampled_test, pred)


precision, recall, _ = precision_recall_curve(y_oversampled_test, pred)

plt.figure(figsize=(15, 15))
plt.step(recall, precision, color='b', alpha=0.1, where='post')
plt.fill_between(recall, precision, step='post', alpha=0.1)
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('Precision-Recall curve: Avg Precision={0:0.2f}'.format(
    average_precision))
```

	precision	recall	f1-score	support
Normal	1.00	1.00	1.00	56750
Fraud	1.00	1.00	1.00	56976
avg / total	1.00	1.00	1.00	113726

Accuracy Score: 0.9998329317834093

Out[20]: Text(0.5,1,'Precision-Recall curve: Avg Precision=1.00')



```
In [25]: def plot_confusion_matrix(cm, classes,
                                     normalize=False,
```

```

        title='Confusion matrix',
        cmap=plt.cm.Blues):

    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')


cnf_matrix = confusion_matrix(y_oversampled_test, pred)
np.set_printoptions(precision=2)

# Plot non-normalized confusion matrix
plt.figure(figsize=(5, 5))
plot_confusion_matrix(cnf_matrix, classes=class_names,
                      title='Confusion matrix, without normalization')

# Plot normalized confusion matrix
plt.figure(figsize=(5, 5))
plot_confusion_matrix(cnf_matrix, classes=class_names, normalize=True,
                      title='Normalized confusion matrix')

plt.show()

```


Confusion matrix, without normalization

```
[[56734  16]  
 [  3 56973]]
```

Normalized confusion matrix

```
[[1.00e+00 2.82e-04]  
 [5.27e-05 1.00e+00]]
```

